

Security in Software

Posix, sanitizers, typical system programming errors, fuzzing

Ole Christian Eidheim

IDI, NTNU

September 28, 2024

Overview

- **POSIX and compilation systems**
- Sanitizers
- Integer overflow
- Range types
- Contracts
- Buffer overflow
- Buffer overread
- Fuzzing
- Exercise 5

Portable Operating System Interface (POSIX)

- IEEE standard for operating systems compatibility
 - most major OSs apart from Windows support this standard
 - POSIX may be too extensive for microkernels:
 - Arduino does not support the POSIX standard
 - Magenta kernel of Google Fuchsia has POSIX compatibility features
 - POSIX contains for example standardized:
 - processes
 - threads
 - (a)synchronous I/O
 - command line interpretation (for instance: `echo "hello" > file.txt`)
 - C libraries
 - Segmentation / Memory Violations
 - command line tools (e.g. `grep`, `kill`, `ls`, `sort`, `tail`, `vi`)
 - Linux is largely POSIX compliant
 - Linux Standard Base: contains extensions that are not part of POSIX
 - MacOS is POSIX certified

POSIX based compiler systems

- collection of compilers, libraries and tools

- GNU Compiler Collection (GCC) (1987)

- gcc/g++

- The most common C/C++ compiler

- Low Level Virtual Machine (LLVM) (2003)

- clang/clang++

- Standard in FreeBSD and MacOS

- Windows supported, [Chrome for Windows](#) kompileres nå med clang++ (2018)

- Used by the Rust compiler

- Various support tools:

- LLDB - less resource-intensive debugger

- Clang Static Analyzer - more thorough static analysis

- Various programs to format/tidy/analyze source code

Overview

- POSIX and compilation systems
- **Sanitizers**
- Integer overflow
- Range types
- Contracts
- Buffer overflow
- Buffer overread
- Fuzzing
- Exercise 5

Sanitizers

- Runtime checks

- Address sanitizer - keeps track of used memory when running the program, and interrupts the program with an error message when an invalid memory operation is performed
 - For example buffer overflow/overread
 - [Valgrind](#) - more comprehensive program for finding such errors
 - Eksempler: [buffer-overflow](#) - demonstrates use of address sanitizer and [Valgrind](#)
- Thread sanitizer - detects *data races*
 - For example, when a variable is altered without using a mutex in two different threads
- Undefined behavior sanitizer - finds for example
 - Signed/unsigned integer overflow
 - Conversion overflow
 - Division by 0
- Both GCC and LLVM contain sanitizers
- Sanitizers are activated solely during testing due to increased resource usage

Overview

- POSIX and compilation systems
- Sanitizers
- **Integer overflow**
- Range types
- Contracts
- Buffer overflow
- Buffer overread
- Fuzzing
- Exercise 5

Integer overflow

- C/C++: prioritize more effective machine code

- Integer overflow is one of many undefined behaviors that must be avoided in critical software

```
#include <iostream>
```

```
int main() {  
    int num = 2147483647;  
    num += 1;  
    std::cout << num << std::endl; // What is the output here?  
}
```

- Some simple integer overflows can be detected at compile time
- Can be detected through [Undefined behaviour sanitizer](#) at runtime

Integer overflow

- Rust: does not accept undefined behaviour

```
fn main() {  
    let mut num: i32 = 2147483647;  
    num += 1; // Program terminates  
    println!("{}", num);  
}
```

- Some simple integer overflows can be detected at compile time
- At runtime, the program terminates if an integer overflow occurs

Integer overflow

- Arbitrary-precision arithmetic

- Languages like Python avoid the problem through number types that can contain numbers of any size (memory is the only limitation)
- Most programming languages have library support for arbitrary-precision arithmetic.
 - In C++: [Boost.Multiprecision](#)
- Resource demanding

Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- **Range types**
- Contracts
- Buffer overflow
- Buffer overread
- Fuzzing
- Exercise 5

Range types in Ada

- Ada: widely used programming language for critical software
 - Used for example in aircraft and flight control software, satellites and banking systems

Range type example in Ada:

```
with Ada.Text_IO; use Ada.Text_IO;

procedure Main is
  type My_Int is range 1 .. 20;
  A : My_Int := 12;
  B : My_Int := 15;

  M : My_Int := (A + B) / 2; -- No overflow here

  N : My_Int := A + B;      -- Will cause compiler warning and
                           -- exception at runtime
begin
  Put_Line ("Hello World!");
end Main;
```

Range types in C++

- `bounded::integer`

- “provides more safety than Ada range types and more efficiency than the C / C++ built-in integer types.”

Example:

```
bounded::integer<0, 10> num(5);  
num+=10; // Compile time error or runtime exception
```

Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- Range types
- **Contracts**
- Buffer overflow
- Buffer overread
- Fuzzing
- Exercise 5

Contracts in Ada

- In Ada, you can define contracts for functions in the form of *pre-* and *post-conditions*.

Example:

```
generic
  type Item is private;
package Stacks is
  type Stack is private;
  function Is_Empty(S: Stack) return Boolean;
  function Is_Full(S: Stack) return Boolean;
  procedure Push(S: in out Stack; X: in Item)
    with
      Pre => not Is_Full(S),
      Post => not Is_Empty(S);
  procedure Pop(S: in out Stack; X: out Item)
    with
      Pre => not Is_Empty(S),
      Post => not Is_Full(S);
private
  ...
end Stacks;
```

Contracts in C++

- Might arrive in the 2026 standard
 - Can be activated during testing, and deactivated in production (for less resource use)
 - Can already be tried out in [Compiler Explorer](#)

Example:

```
int f(int n)
[[pre: n >= 0]]
[[post r: n == r]]
{
    return n;
}

int main() {
    f(-2); // Contract violation
}
```


Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- Range types
- Contracts
- **Buffer overflow**
- Buffer overread
- Fuzzing
- Exercise 5

Buffer overflow

```
char          A[8] = {};  
unsigned short B = 1979;
```

variable name	A								B	
value	[null string]								1979	
hex value	00	00	00	00	00	00	00	00	07	BB

```
strcpy(A, "excessive");
```

variable name	A								B	
value	'e'	'x'	'c'	'e'	's'	's'	'i'	'v'	25856	
hex	65	78	63	65	73	73	69	76	65	00

Buffer overflow

- editing data, the stack, example 1

```
#include <iostream>

using namespace std;

int main() {
    char chr = 'a';
    char table[3] = {1, 2, 3};

    cout << "table starts at memory address "
          << &table << endl;
    cout << "chr starts at memory address "
          << (void *)&chr << endl;

    table[3] = 'b';

    cout << chr << endl;
}
```

Output e.g:

table starts at memory address 0x7fff51ed6b98

chr starts at memory address 0x7fff51ed6b9b

b

- Protection of objects in the stack
 - Use high-level programming if possible
 - Foreach loops or **functional algorithms**
 - Use support tools such as:
 - AddressSanitizer
 - Valgrind
 - The compiler may issue warnings/error messages

Buffer overflow

- editing data, the stack, example 2

```
#include <stdio.h>
#include <string.h>

char input[20];
int success;

int main() {
    success = 0;

    printf("Password: ");           // Print to standard output
    fflush(stdout);                 // Flush standard output
    scanf("%s", input);             // Read from standard input,
                                    // and store the data in the variable input
    if (strcmp(input, "PassWord213") == 0) // If input is equal to: PassWord213
        success = 1;

    if (success)
        printf("Logged in\n");
    else
        printf("Wrong password\n");
}
```

Buffer overflow

- editing data, the heap

```
#include <iostream>
#include <vector>

using namespace std;

int main() {
    vector<char> table = {1, 2, 3};
    char *chr = new char('a');

    cout << "table.data() starts at memory address "
          << (void *)table.data() << endl;
    cout << "chr starts at memory address "
          << (void *)chr << endl;

    table[16] = 'b';

    cout << *chr << endl;
}
```

Output e.g.:

table.data() starts at memory address 0x7f9929500000

chr starts at memory address 0x7f9929500010

b

- Protection of data in the heap
 - Use high-level programming if possible
 - Foreach loops or [functional algorithms](#)
 - Use support tools such as:
 - AddressSanitizer
 - Valgrind
 - Range check with exception possible
 - e.g. table.at(16)

Buffer overflow

- editing function pointers

IKKE KJØR DENNE KILDEKODEN

```
#include <iostream>

using namespace std;

void func1() {
    cout << "Output from func1" << endl;
}

void func2() {
    cout << "Output from func2" << endl;
}

int main() {
    cout << "The function func1 starts at memory address " << (void *)&func1 << endl;
    cout << "The function func2 starts at memory address " << (void *)&func2 << endl;

    auto f = &func1;
    char table[3] = {1, 2, 3};

    cout << "The pointer f starts at memory address " << (void *)&f << endl;
    cout << "The table starts at memory address " << (void *)&table << endl;

    f();

    table[3] = (size_t)&func1 + 64;

    f();
}
```

Output e.g.:

The function func1 starts at memory address 0x10076ff90

The function func2 starts at memory address 0x10076ffd0

The pointer f starts at memory address 0x7fff5f490b80

The table starts at memory address 0x7fff5f490b7d

Output from func1

Output from func2

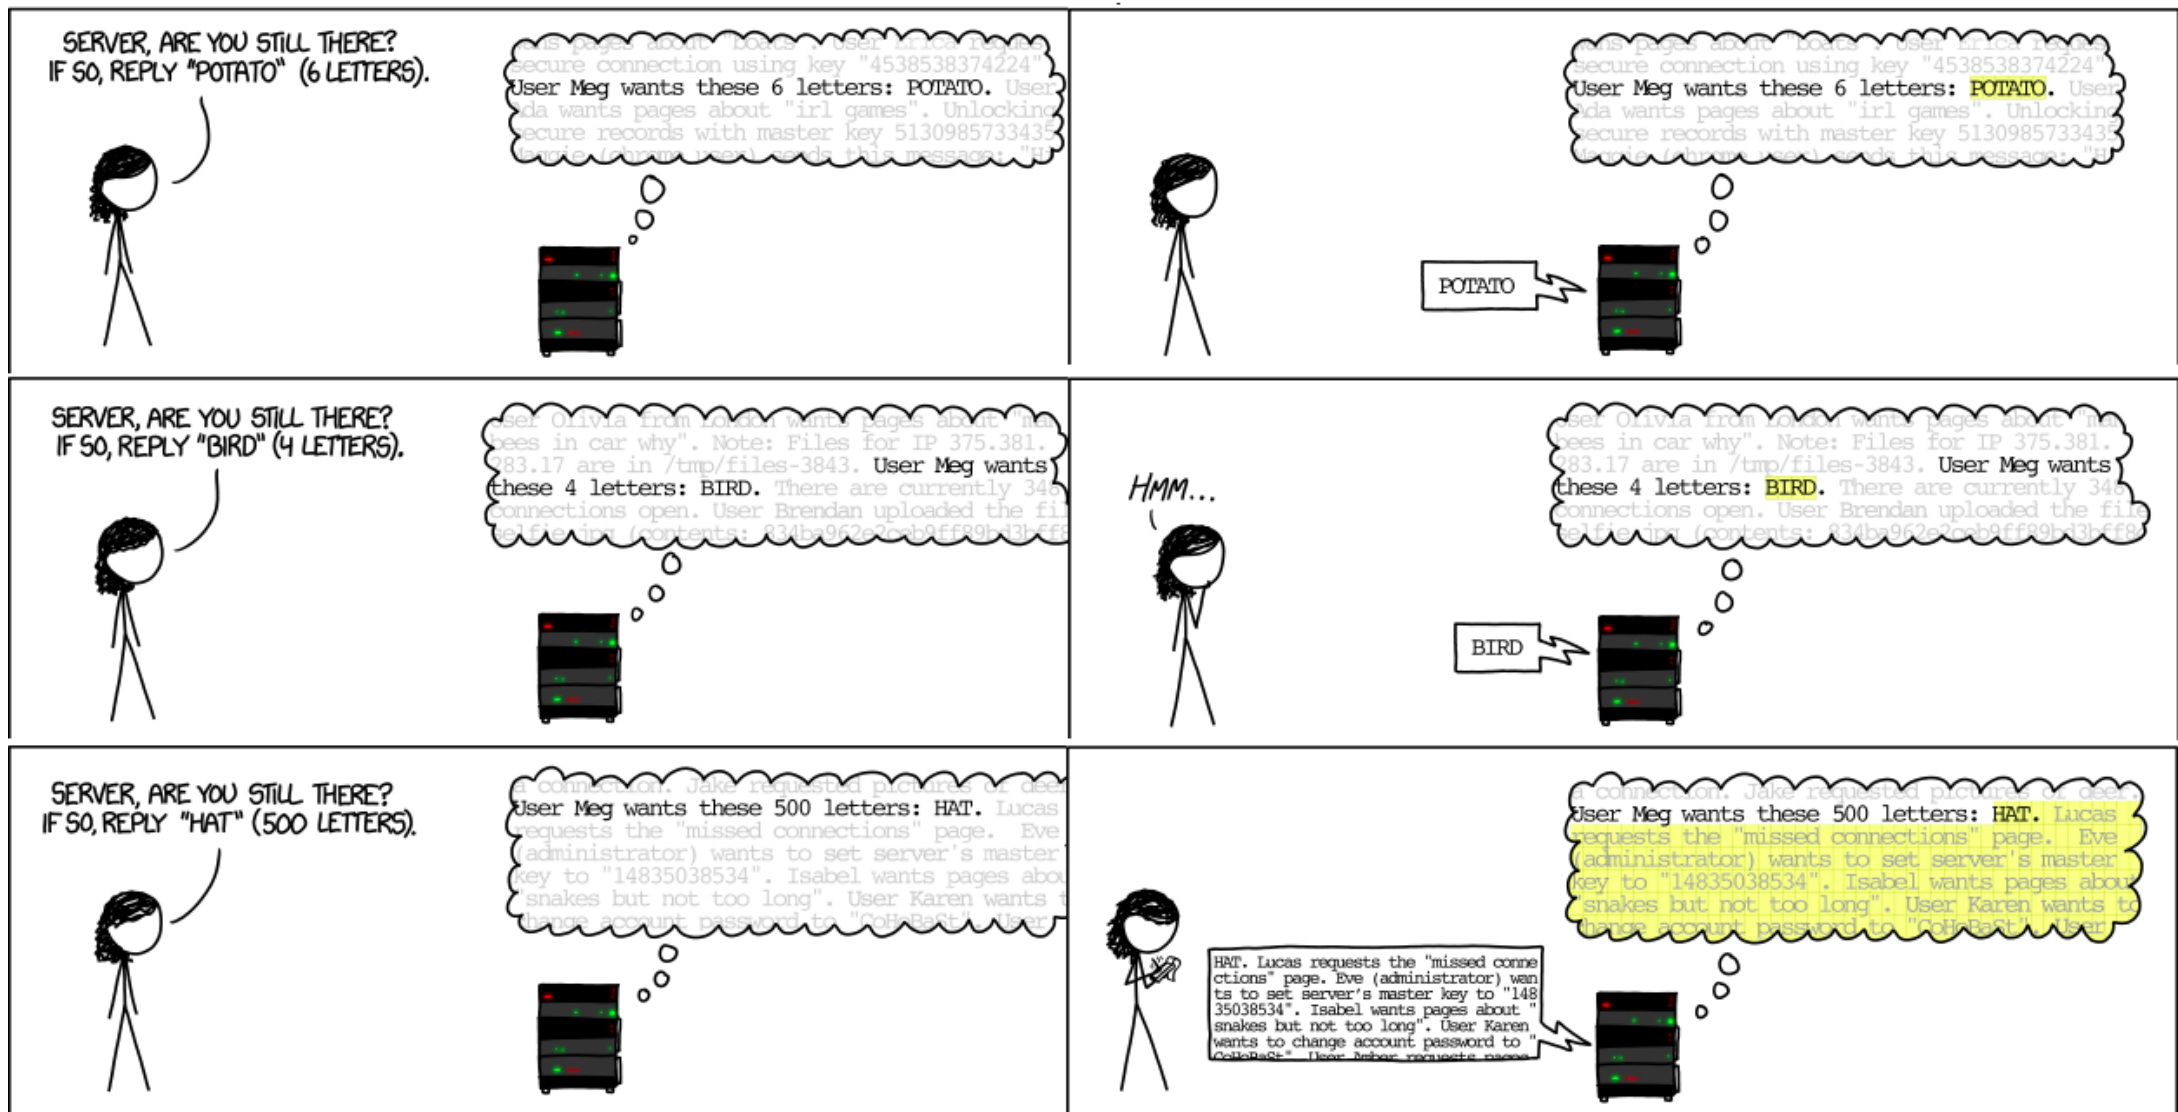
Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- Range types
- Contracts
- Buffer overflow
- **Buffer overread**
- Fuzzing
- Exercise 5

Buffer overread

- Heartbleed (OpenSSL vulnerability for versions <1.0.1g)

HOW THE HEARTBLEED BUG WORKS:



Buffer overread

- Heartbleed-like example

```
#include <iostream>
#include <string>

using namespace std;

int main() {
    string secret_data="secret key!";

    string str;
    int length;
    cin >> str >> length;

    for(int c = 0; c < length; c++)
        cout << str[c];

    cout << endl;
}
```

Input:

hello 40

Output e.g.:

hello???????????????????? secret key!

Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- Range types
- Contracts
- Buffer overflow
- Buffer overread
- **Fuzzing**
- Exercise 5

Fuzzing

- Automatic testing of functions through generated input parameters
 - Popular tool: [libFuzzer](#)
 - Have been used to find [bugs in widely used software](#)
 - Works together with sanitizers
 - Advanced functionality: “... the fuzzer then tracks which areas of the code are reached, and generates mutations on the corpus of input data in order to maximize the code coverage.”
- See [fuzzing-example](#)

Overview

- POSIX and compilation systems
- Sanitizers
- Integer overflow
- Range types
- Contracts
- Buffer overflow
- Buffer overread
- Fuzzing
- **Exercise 5**

Exercise 5

- Perform *fuzzing* with *address sanitizer* on the C function you created in exercise 4b)
 - Fix bugs you find through fuzzing, or introduce bugs that are discovered through fuzzing
 - **Voluntary:** Set up a CI (Continuous Integration) solution that performs *fuzzing* with *address sanitizer*
 - See [instructions for running fuzzing in a terminal](#) (note that you can limit the number of seconds libFuzzer will run with the `-max_total_time` argument)
 - Create tests, and run the tests as well in CI ([fuzzing-example](#) contains an example test)